The basic: Algorithm Implemented.

Answer: we update the CartPole and Lunar Lander environment using the Double DQN Algorithm.

The Double DQN algorithm implemented uses an Experience Replay, which is used to increase the effectiveness and stability of learning. This involves randomly selecting a batch of experiences from the replay buffer to train the agent's neural network after storing the agent's experiences (i.e., observations, actions, rewards, and next states) in the replay buffer. Randomly sampling a batch of experiences from the replay buffer, the agent's neural network can break the correlation between consecutive experiences, which can improve learning efficiency and stability.

We implemented Double DQN using torch library on the ‘CartPole-v1’, The Lunar Lander V-2 environment, and the Grid-world environment. The target network is used for calculating the Q-value of the action selected by the primary network in the next state. This will solve the overestimation problem in calculating Q-values that can occur in standard DQN implementation in RL environments.

improvement over the vanilla DQN

Double DQN addresses the issue of overestimation of Q Values which results in more precise value estimates and better performance.
When compared to DQN, double DQN is more trainable and less sensitive to the selection of hyperparameters with a cost of more computational requirements. The Double DQN uses two separate networks to select and evaluate the action and update the Q network's weights. We used the Neural Network that we built using the Pytorch libraries to implement the improved version of vanilla DQN.

Double DQN

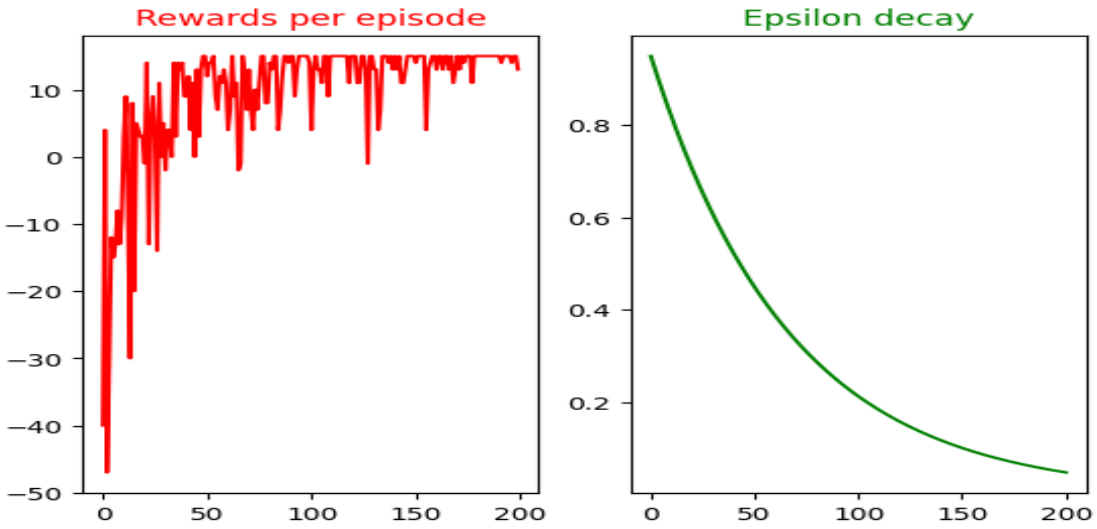

CartPole-v1:

Double DQN:

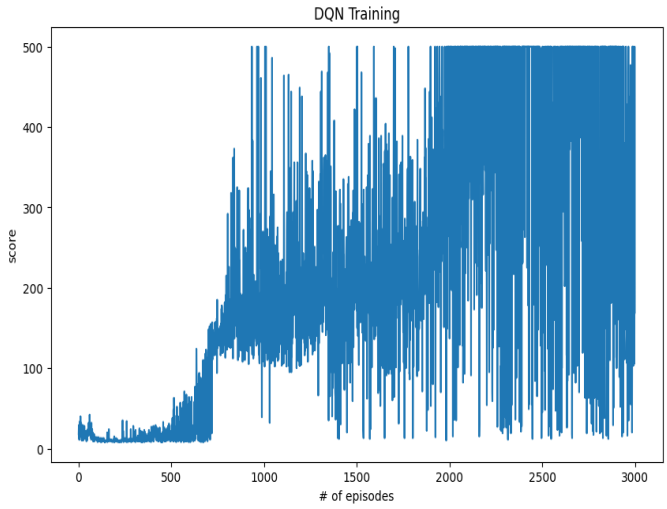

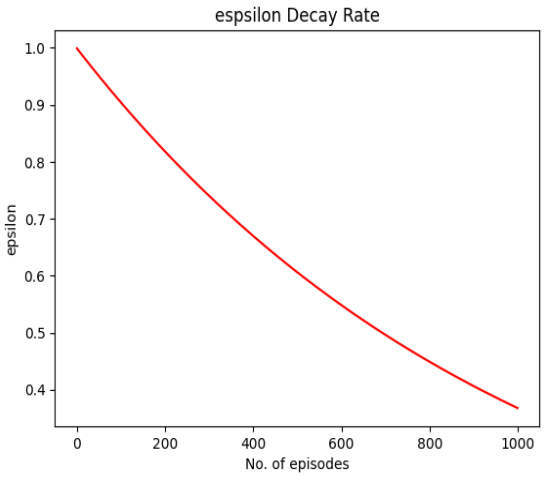

Lunar Lander -V2:

Double-DQN: The lunar lander is a box-2d environment, Where the agent, The lander must land on the moon. It's a complex environment to solve with higher state and action space. The D-Dqn was able to solve the environment and we can observe the the lander achieving a 200+ rewards and being able to maintain the moving average. The results are plotted below as reward per episode and epsilon decay.

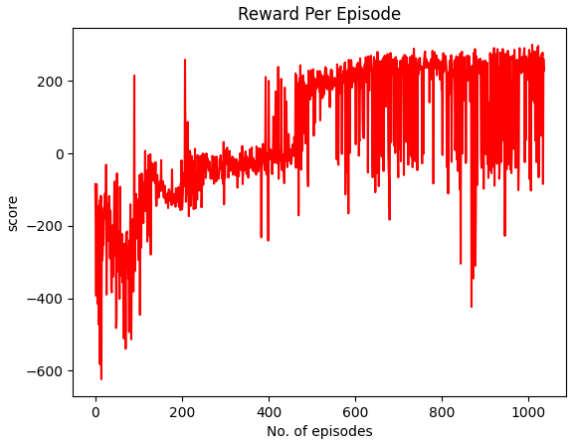

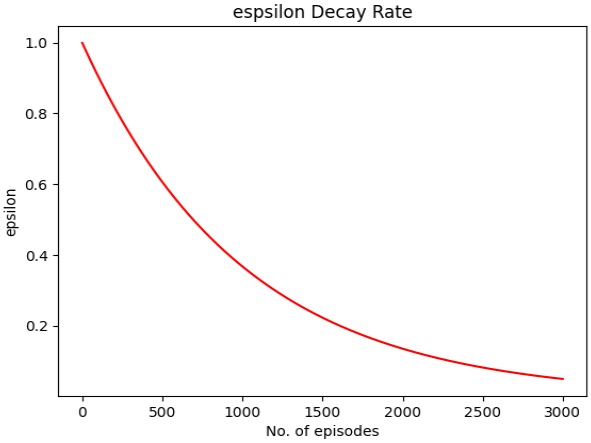

environment run for 5 episodes, *where the agent chooses only greedy actions from the learned policy.*

GridWorld

Double DQN

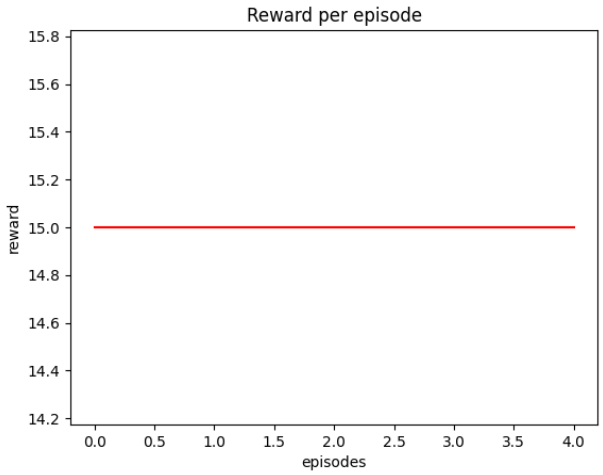

DQN

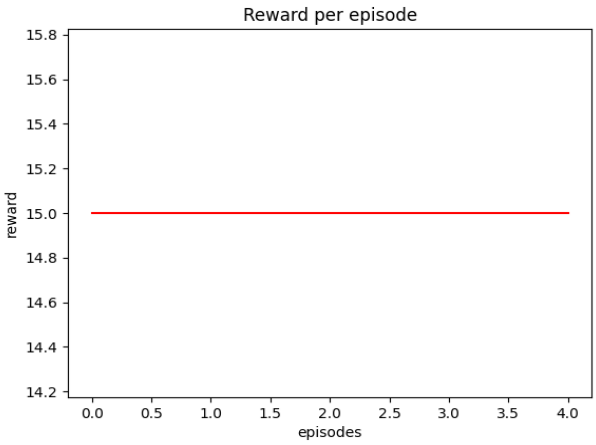

CartPole-v1

Double DQN

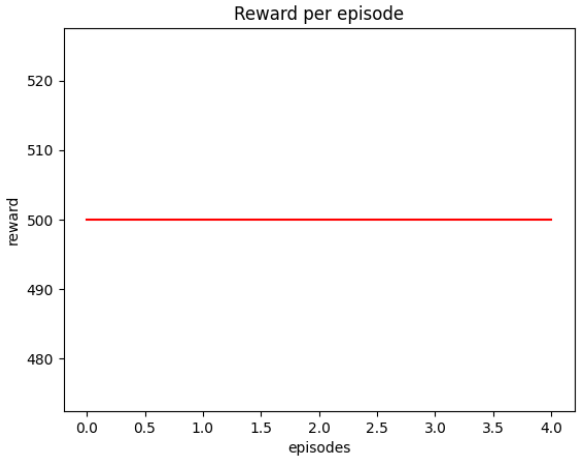

DQN

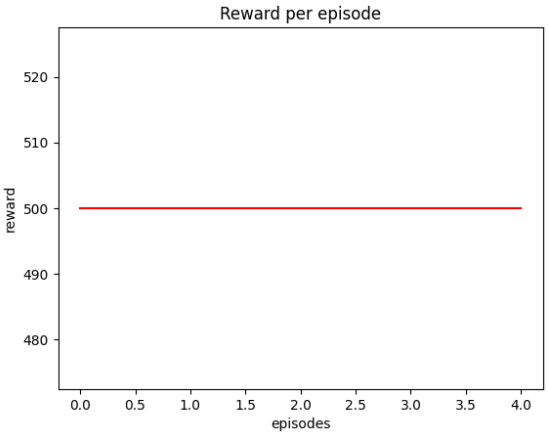

LunarLander-v2

Double DQN

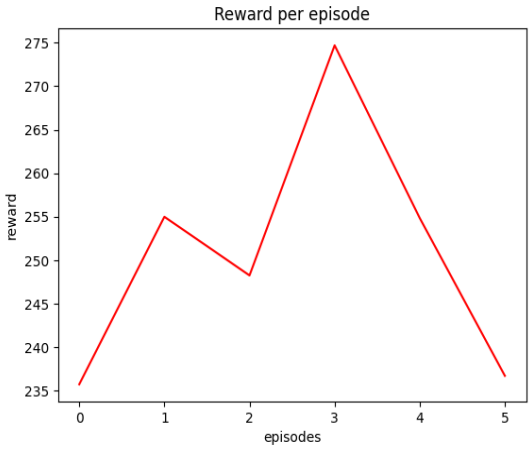

DQN

performance of both algorithms (DQN & doubleDQN) on *the same environments with three rewards dynamics plots with results from two algorithms applied on:*

- *Grid-world environment*
- *'CartPole-v1'*
- *'LunarLander-v2'*

Both algorithms achieved the optimal rewards in the defined grid world environment. The Double DQN took the path of exploring more and gradually started to exploit the q-network. Dqn followed a similar path in this environment.

Double DQN

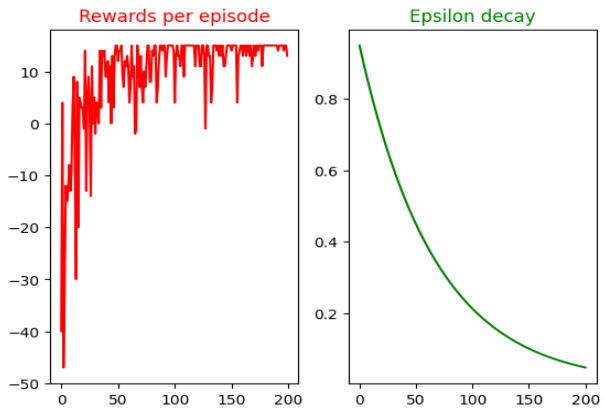

DQN

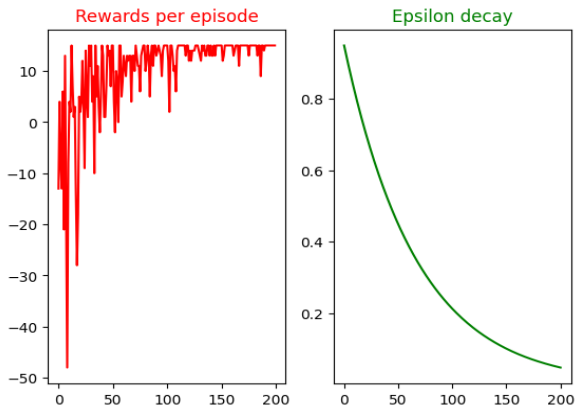

CartPole-v1:

The DQN started with more negative regard and slowed learning the optimal policy from the network while the Double DQn used the exploration more for the first many episodes and slowly used the q network to achieve the convergence and solve the environment. Overall Bothe algorithms were able to solve.

DQN

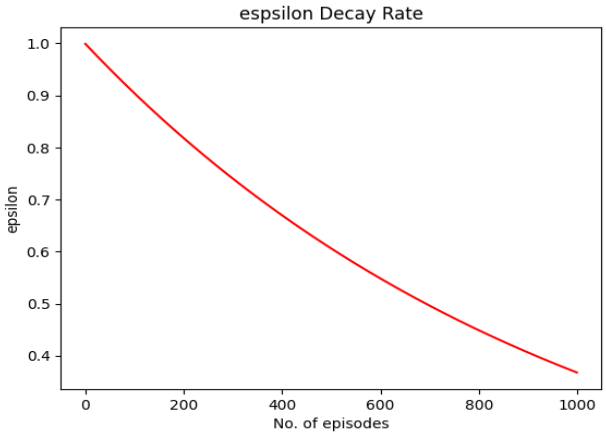

Double DQN:

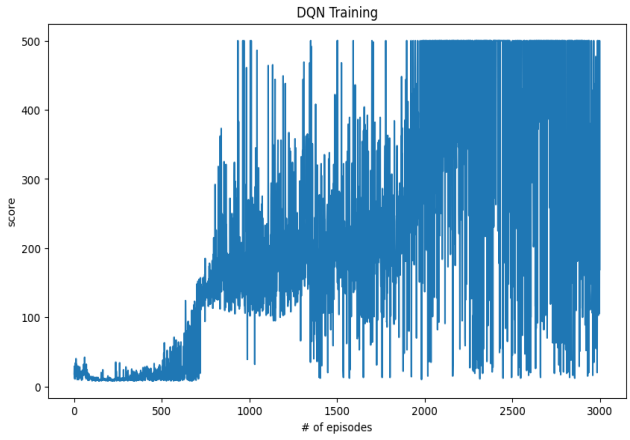

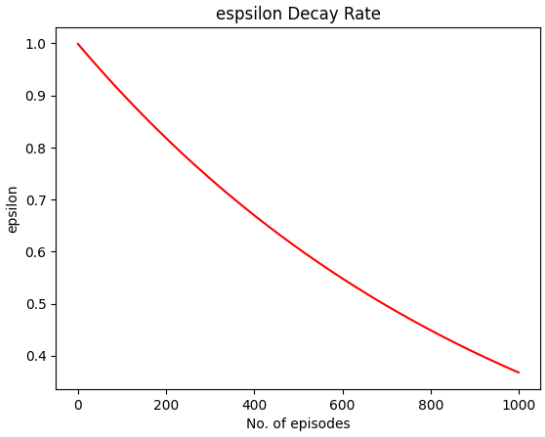

Lunar Lander -V2:

We have solved the Lunar Lander environment, using Deep Q-Network (DQN). However, a Double Q-Network (DQN) that solves the overestimation issue in Q-learning can be used to further boost its performance. The Double DQN was able to solve the environment and was able to learn the optical path.

Double-DQN

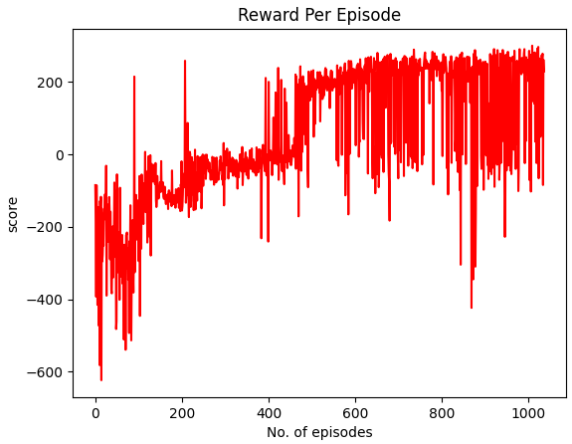

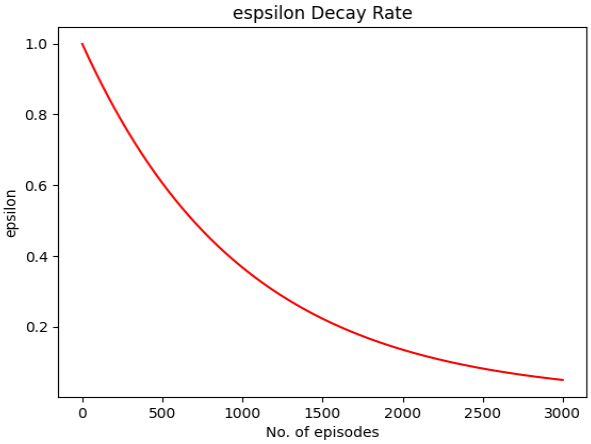

DQN

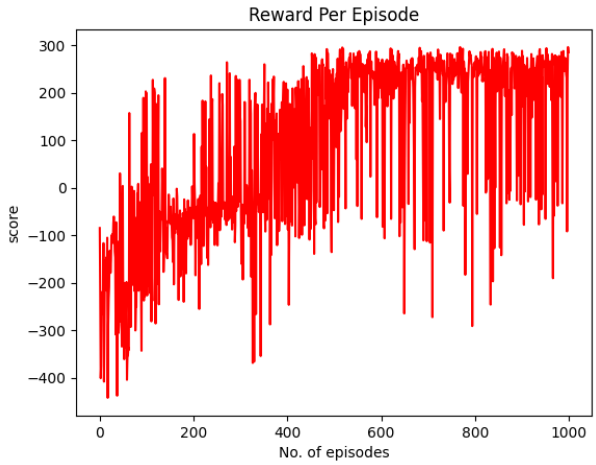

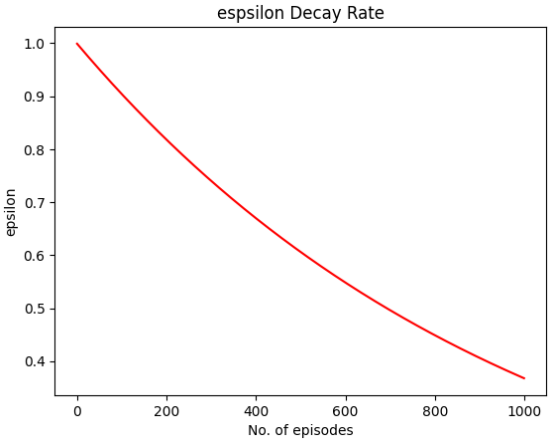

My views: same algorithm different birds*, or how various algorithms behave in the same environment*

Depending on the env and the parameter that the Algorithms is applied on, The performance of a reinforcement learning algorithm like DQN or Double DQN can change. The Double DQN solves the overestimation problems but can be more expensive to train since it uses two different networks, which requires more processing and memory than DQN. It took 10-mins of Google's standard TPU time to just train the network.

While DQN is simpler to train and can quickly converge in simpler environments like grid or cart pole environments. The overestimation problem cause the agent to take no optimal actions. The DQN is also very sensitive to hyperparameters.

RL Part-2
Assignment 2

Name 1) Mantri, Sai Rochan
Ubit: vmantri
Email:vmantri@buffalo.edu

Name 2) Vanam, Saishankar
Ubit: saishank
Email:saishank@buffalo.edu

Team Member	Assignment Part	Contribution (%)
Mantri, Sai Rochan	II	50%
Vannam, Saishankar	II	50%

1. Discuss the benefits of:
    - Using experience replay in DQN and how its size can influence the results
    - Introducing the target network
    - Representing the Q function as qˆ(s, w)

Ans:

- Using experience replay in DQN and how its size can influence the results

In Deep Q-networks (DQNs), Experience Replay is used to increase the effectiveness and stability of learning. This involves randomly selecting a batch of experiences from the replay buffer to train the agent's neural network after storing the agent's experiences (i.e., observations, actions, rewards, and next states) in the replay buffer. Randomly sampling a batch of experiences from the replay buffer, the agent's neural network can break the correlation between consecutive experiences, which can improve learning efficiency and stability.

In this Assignemnt, We implemented DQN using torch librarry on the ‘CartPole-v1’, the second complex environment, and the grid-world environment. The DQN uses a replay buffer to push the agent’s experiences (shown in below) from which a batch is sampled out and used to train the agent’s neural network. The results looked much better when a using a buffer and using the experience replay to train the DQN.

```
print(showASampleBuffer)
```

```
{'state': array([0.  , 0.  , 0.  , 0.  , 0.  , 0.  , 0.23, 0.  , 0.  , 0.22, 0.24,
       0.  , 0.  , 0.  , 0.5 , 0.  ]), 'action': 1, 'reward': 10, 'next_state': array([0.  , 0.  , 0.  , 0.  , 0.  , 0.  ,
0.23, 0.  , 0.  , 0.22, 0.24,
       0.  , 0.  , 0.  , 0.5 , 0.  ]), 'done': False}
```

• Introducing the target network

In deep Q-learning, introducing a target network is a technique used to improve the stability and efficiency of learning. The target network is a separate neural network that is used to estimate the Q-values of the next state in the Q-learning update, instead of using the same network that is being updated during training.

In Impelemented Models, We used a target network and policy network, Two seperate instances of the DQN. The primary network is used to select actions that maximize the expected future reward, while the target network is used to estimate the Q-values of the next state.

```
print("action from Traget Network :",target_net.sample_action(torch.from_numpy(state).float()).
```

```
action from Traget Network : 1
```

```
print("action from Primary Network :",prime_net.sample_action(torch.from_numpy(state).float()).
```

```
action from Primary Network : 1
```

• Representing the Q function as qˆ(s, w)

The Q-function (or state-action value function) is a function that maps a state-action pair to the expected cumulative reward. Representing the Q-function in high-dimensional state and action spaces in complex environments with complex state and action spaces. This makes it difficult to represent the Q-function explicitly using a lookup table or other conventional methods. This were we used a neural network for function approximation.

Q-function is represented as a parameterized function qˆ(s, w), where s is the state, w is the weight vector of the function, and qˆ(s, w) is the estimated value of the state-action

pair. We used this parameterized Q-function qˆ(s, w) as a instances of the DQN class in the implemented environments.
2. Briefly describe ‘CartPole-v1’, the second complex environment, and the grid-world environment that you used (e.g. possible actions, states, agent, goal, rewards, etc). You can reuse related parts from your Assignment 1 report.

Grid Environment:

Gridworld is a simple environment that can be used for testing different algorithms that we developed as part of the assignment. The environment consists of a two-dimensional grid with cells, where each cell can be either empty, an reward/monster and a goal state or terminal state. The agent is located in a specific cell (0,0) in the grid world we defined and can take 4 possible actions at particular state, move up, down, left or right to neighboring cells. The goal is to reach the goal cell from the starting cell while collecting the highest reward possible.

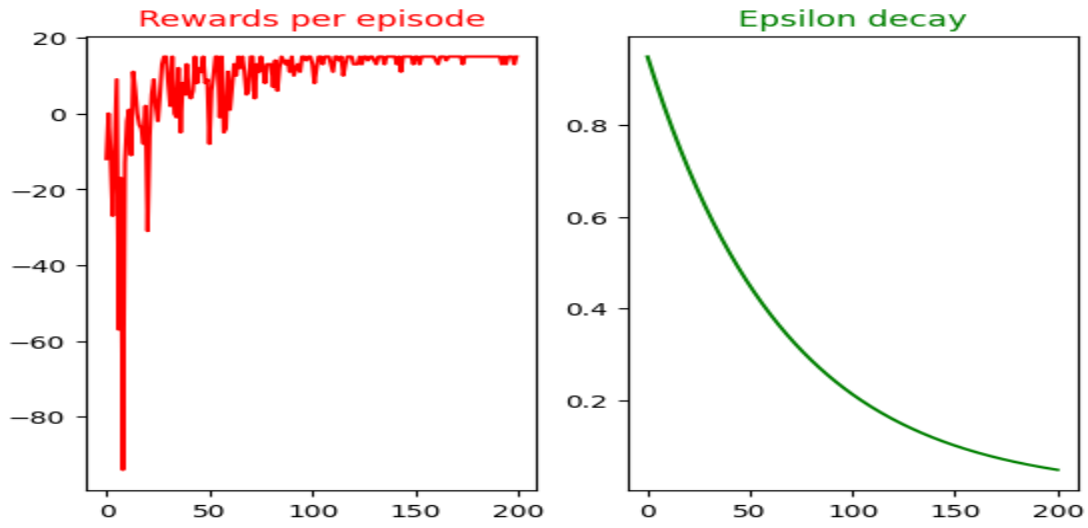

The state of the environment is represented by the position of the agent in the grid, and the actions are the four possible directions the agent can move. The reward for moving into a cell is -1, except for when the agent reaches the goal, in which case the reward is +10. If the agent tries to move into one of these rewards states [1,2], [2,1] and [2,2], The agent bags a reward. The [2,1] state has a monster with a reward of negative 3.

The Gridworld environment can be configured to have different sizes, shapes, and numbers of rewards/monsters and goals, making it a good testing environment for the implemented DQN Algorithm using torch libraries. We used the DQN on the grid world

enovirmnt and plotted the results for Reward per Episode and Epsilon Decay as shown above. The agent was able to get the highest reward with trained DQN's help.

CartPole-v1

The main goal of the environment is to balance a pole on top of a cart by moving the cart left or right. This solved state is when the pole remains upright. The environment is considered "Done" if the pole remains upright for a specified period of time or the pole have moved to an angle out of reach.

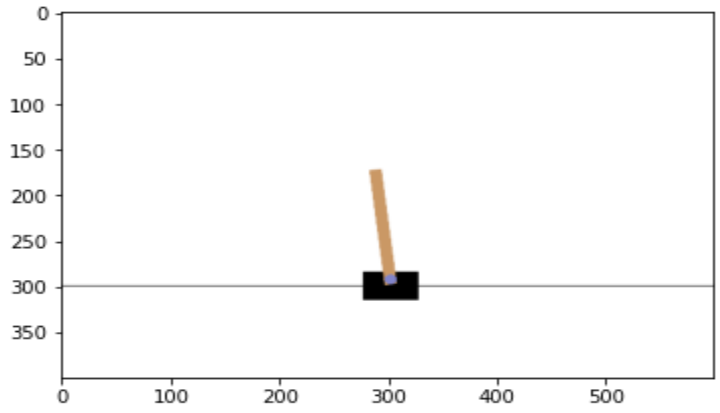

[The position and velocity of the cart, The angle and angular velocity of the pole] make the states of the Cart Pole Environemnt, and actions are to move the cart either to the left or the right (binary). For each time step that the pole remains upright, a reward of +1 is added to the list to plot the graphs. A negative is awarded if game is over.

We have used the Implemented DQN on the Cartpole-v1 environment and the DQN is able to successfully solve the environment as defined by the gym as the average reward over number of episodes. The results are plotted against the Rewards per Episode and Episoln decay as shown below.

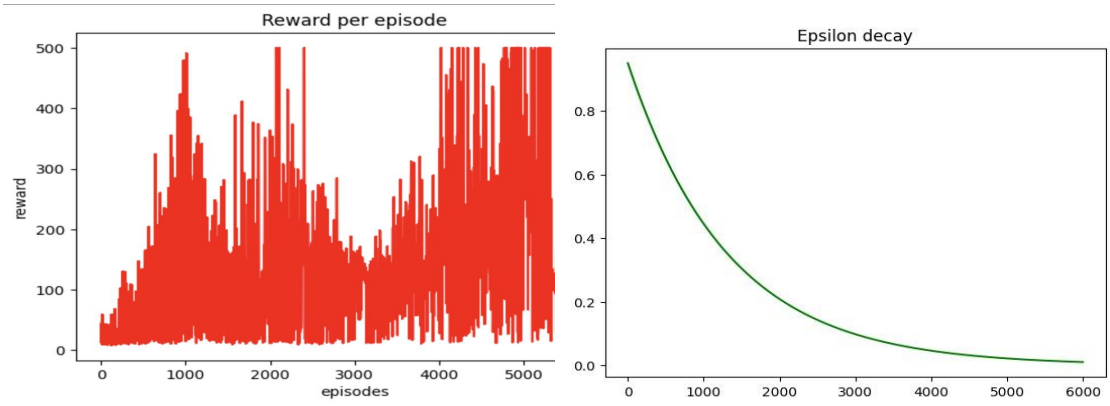

Lunar Lander

The Moon Landing Environment is a somewhat complex OpenAI gym environment. An agent (a spacecraft) must learn how to safely land on the moon's surface in an environment that resembles a lunar landing situation. The environment allows the agent to manage the spacecraft's main engine, side engines, and thrusters as well as receive sensor readings including altitude, speed, and direction.

The agent, in our case- A Lunar Lander, in the Lunar Landing Environment has actions as a 2-dimensional vector representing the amount of thrust to apply to the main engine and side engines. The goal of the enviornemnt is to get 200 points. We have the trained DQN Agent on Lunar Landing Enviroment and plotted Rewards per Episode and Epsilon decay.

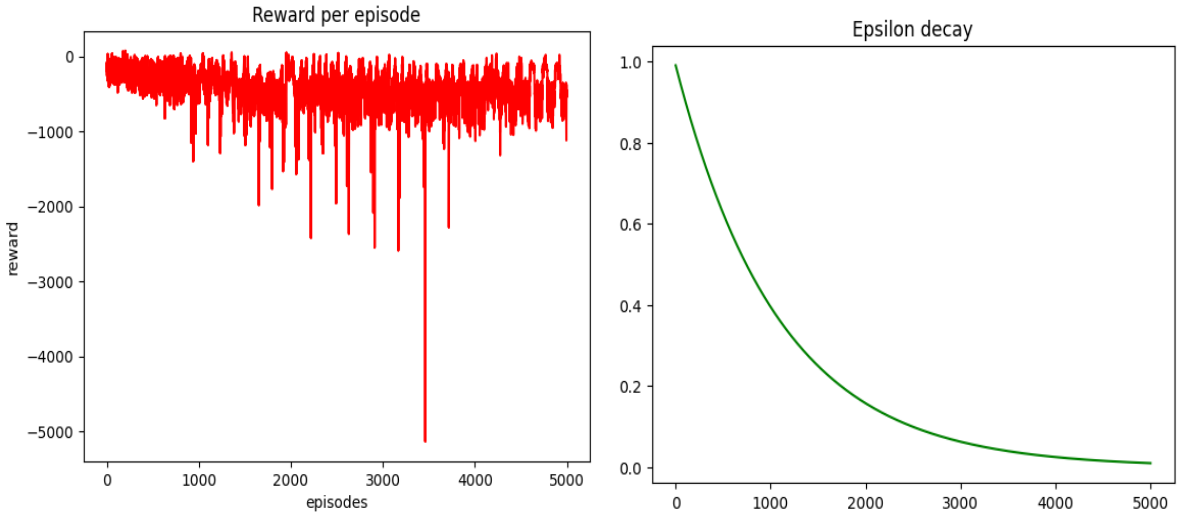

4. Provide the evaluation results. Run your agent on the three environments for at least 5 episodes, where the agent chooses only greedy actions from the learnt policy. Plot should include the total reward per episode.

1) GridWorld environemnt

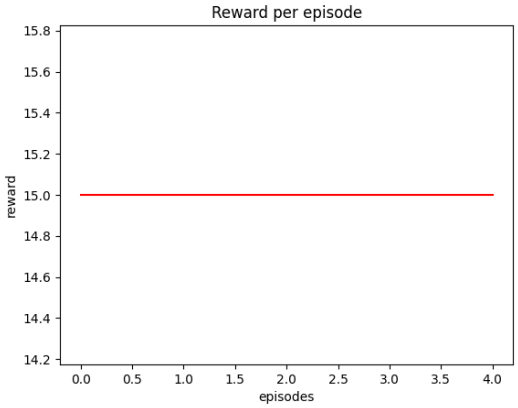

2) Cart-Pole

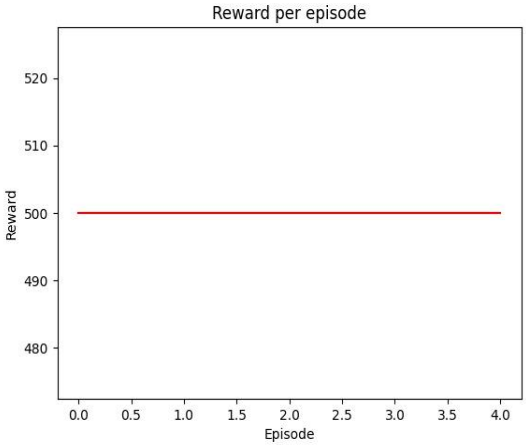

3) Lunar Landing

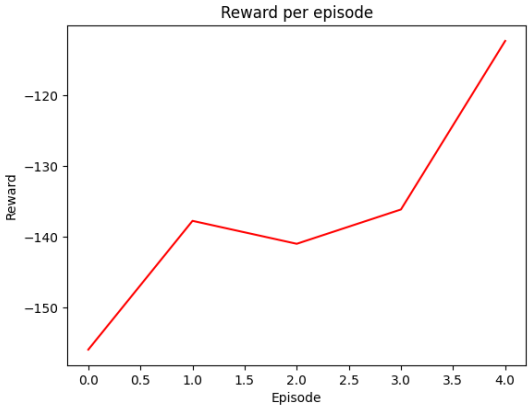

5. Provide your interpretation of the results. E.g. how the DQN algorithm behaves on different envs.

DQN algorithm is an effective deep reinforcement learning technique that can be used to train agents in environments with large state spaces. As Discuused above, We have used a DQN to train the agent and created Instaces of DQN class to introduce a target network and have a replay buffer for the Agnet to learn from as a backup.

The cart pole and Lunar Landing environment has a large and continuous state space, which makes it impractical to define a Q-table for each possible state-action pair. Instead, a function approximator is used to estimate the Q-values for each state-action pair, and the DQN algorithm is employed to train the agent to maximize its reward. The loss function is used, which quantifies the difference between the predicted and actual Q-values. The Adam optimizer is used to adjust the weights of the neural network, which helps improve the accuracy of the Q-value estimates and the overall performance of the agent.

In LunarLander-v2, The agent observes the environment in the form of a two-dimensional image, and the DQN network learns to map images to actions that maximize the expected cumulative reward. However, LunarLander-v2 is a more complex environment than CartPole-v1, and the performance of the DQN algorithm was very highly dependent on hyperparameters. We Noticed are the size of the batch and exploration-exploitaion techniques and size of nueral network (DQN) itself.

References:

https://www.gymlibrary.dev/environments/classic_control/cart_pole/
https://www.gymlibrary.dev/environments/box2d/lunar_lander/
Prof. Alina Vereshchaka. Lecture Slides.

https://web.stanford.edu/class/psych209/Readings/MnihEtAlHassibis15NatureControlDeepRL.pdf
https://www.youtube.com/watch?v=IHZwWFHWa-w&ab_channel=3Blue1Brown
https://www.youtube.com/watch?v=Ilg3gGewQ5U&t=206s&ab_channel=3Blue1Brown